

Intelligent Data Analyst Agent Architecture (MCP-Based)

Objective

To design an intelligent, modular, data-agnostic agent capable of processing diverse traffic-related data using a dynamic suite of up to 100 MCP (Model Context Protocol) tools. The system supports high-throughput workloads, tool chaining, visualization, and secure API-based integration into larger architectures.



Core Components

1. Agent API Layer

- **Tech:** Python (FastAPI/Flask), Docker
 - **Purpose:** Exposes REST endpoints (e.g., `/analyze`) to accept input data + context
 - **Responsibilities:**
 - Parse & validate request payloads
 - Forward to Router for tool assignment
 - Manage response formatting, chaining, and error handling
 - Emit job/run events for visualization dashboards
-

2. Router Module

- **Modes:**
 - **Rule-Based Routing:** Based on input types, keywords, metadata
 - **ML-Based Routing:** Intent classification using ML/LLM
 - **Hybrid Routing:** Fast rules + fallback to model-driven dispatch
 - **Responsibilities:**
 - Decide best-fit MCP tool or tool chain
 - Send routing metadata to Dispatcher
 - Generate trace identifiers and step metadata for visualization
-

3. Tool Dispatcher

- **Function:** Orchestrates tool invocations
- **Protocols Supported:**
 - **REST** (default)
 - **gRPC** (for low-latency/high-throughput)
 - **Kafka/RabbitMQ** (for async workloads)

- **Responsibilities:**
 - Handles retries, timeouts
 - Resolves tool endpoint from Tool Registry
 - Loads authentication headers + payloads
 - Publishes step progress and results for visualization and tracing
-

4. Tool Chaining Manager

- **Purpose:** Executes tool pipelines
 - **Approach:**
 - DAG-based chaining (e.g., Anomaly Detection -> Clustering)
 - Agent-guided dynamic chaining
 - **Responsibilities:**
 - Manage data hand-off between tools
 - Track intermediate results and state
 - Report stage transitions to the visualization subsystem
-

5. Tool Registry

- **Storage:** Local JSON file or database
 - **Fields:**
 - Tool name, task type, supported data types
 - Endpoint URL & communication protocol
 - Version, metadata, health status
 - Visualization metadata: category, icon, color code, owner
-

6. MCP Tool Interface

- **Standardized Schema:**

```
{
  "input": { ... },
  "context": { ... }
}
```

- **Returns:**

```
{
  "status": "success",
  "output": { ... },
  "meta": { ... }
}
```

- **Implementation:**
 - Python class wrapper for tool logic
 - Exposed via REST/gRPC/Kafka listener
 - Emits structured logs, spans, and step metrics for visualization
-

7. Visualization & Monitoring Layer

- **Purpose:** Provide full transparency into the running processes and tool interactions.
 - **Components:**
 - **Run Service / Jobs API:** Tracks all runs, jobs, steps, and statuses; exposes REST + WebSocket endpoints (e.g., `/v1/runs`, `/ws/runs/{id}`).
 - **Dashboard UI:** Shows tool catalog, live runs, DAG visualizer (tool chaining), job progress, and metrics.
 - **Tracing:** OpenTelemetry + Jaeger/Tempo for distributed traces.
 - **Logs:** ELK/OpenSearch for structured logs (linked to runs).
 - **Metrics:** Prometheus + Grafana for latency, throughput, error rate.
 - **Lineage:** OpenLineage/Marquez integration for dataset-tool-output provenance.
 - **User View:**
 - Tool catalog with capabilities, schema, and status.
 - Real-time run status, progress bars, ETA, per-step logs, and trace links.
 - DAG view showing current pipeline execution flow.
-

8. Communication & Throughput Management

- **Options:**
 - REST (development, small-scale)
 - gRPC (binary RPCs, high throughput)
 - Kafka (buffered, async tasks)
 - **Concurrency:**
 - Python asyncio / Celery for parallel calls
 - K8s for container auto-scaling
 - KEDA for queue-based scaling
-

9. Security Layer

- **Auth:**
- JWT-based access control
- HMAC signing for internal tool calls
- **Transport:**
- TLS encryption for REST/gRPC
- Kafka: TLS + SASL
- **Audit Logging:**
- Request, tool, user, timestamp, result status
- Integrated into visualization UI for admin access

10. Observability & Governance

- **Tracing:** OpenTelemetry spans per request and tool
 - **Metrics:** Prometheus collectors for latency, throughput, queue lag
 - **Logging:** Structured, tenant-aware JSON logs
 - **SLOs & Alerts:** Alertmanager for anomalies and health checks
 - **Governance:**
 - Versioning, tool lifecycle tracking
 - Canary releases and shadow runs
 - UI displays deprecation notices and tool change logs
-



Data & Task Support

Input Types:

- Tabular (CSV, Excel, SQL result)
- Text (incident reports, logs)
- JSON/XML (API or IoT device input)
- Images (traffic cams)
- Geo (GeoJSON, GPS points)

Supported Tasks:

- Anomaly Detection
 - Incident Detection
 - Time-Series Forecasting
 - Descriptive Stats & Comparison
 - Classification / Regression
 - Clustering & Feature Engineering
 - Geospatial Mapping & Analysis
-

Deployment Notes

- Containerized with Docker for each tool and core module
 - Use Docker Compose or Kubernetes for orchestration
 - Includes visualization stack (Grafana, Jaeger, ELK) and UI dashboard
 - Designed to plug into larger systems as a callable API service
 - Future-proofed for more advanced ML planning agents (e.g., LLM planner)
-

Next Steps

1. Scaffold base API + agent logic

2. Implement router (rule-based first)
3. Add 3–5 MCP tools with REST endpoints
4. Integrate Run Service + WebSocket for real-time progress
5. Add UI layer for visualization (DAG, runs, logs)
6. Package with Docker Compose for local testing
7. Extend with gRPC + Kafka for async cases